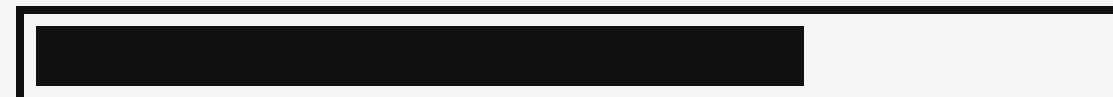




# Event Logging story



*Disclaimer : I am not a Fluent Advocate*



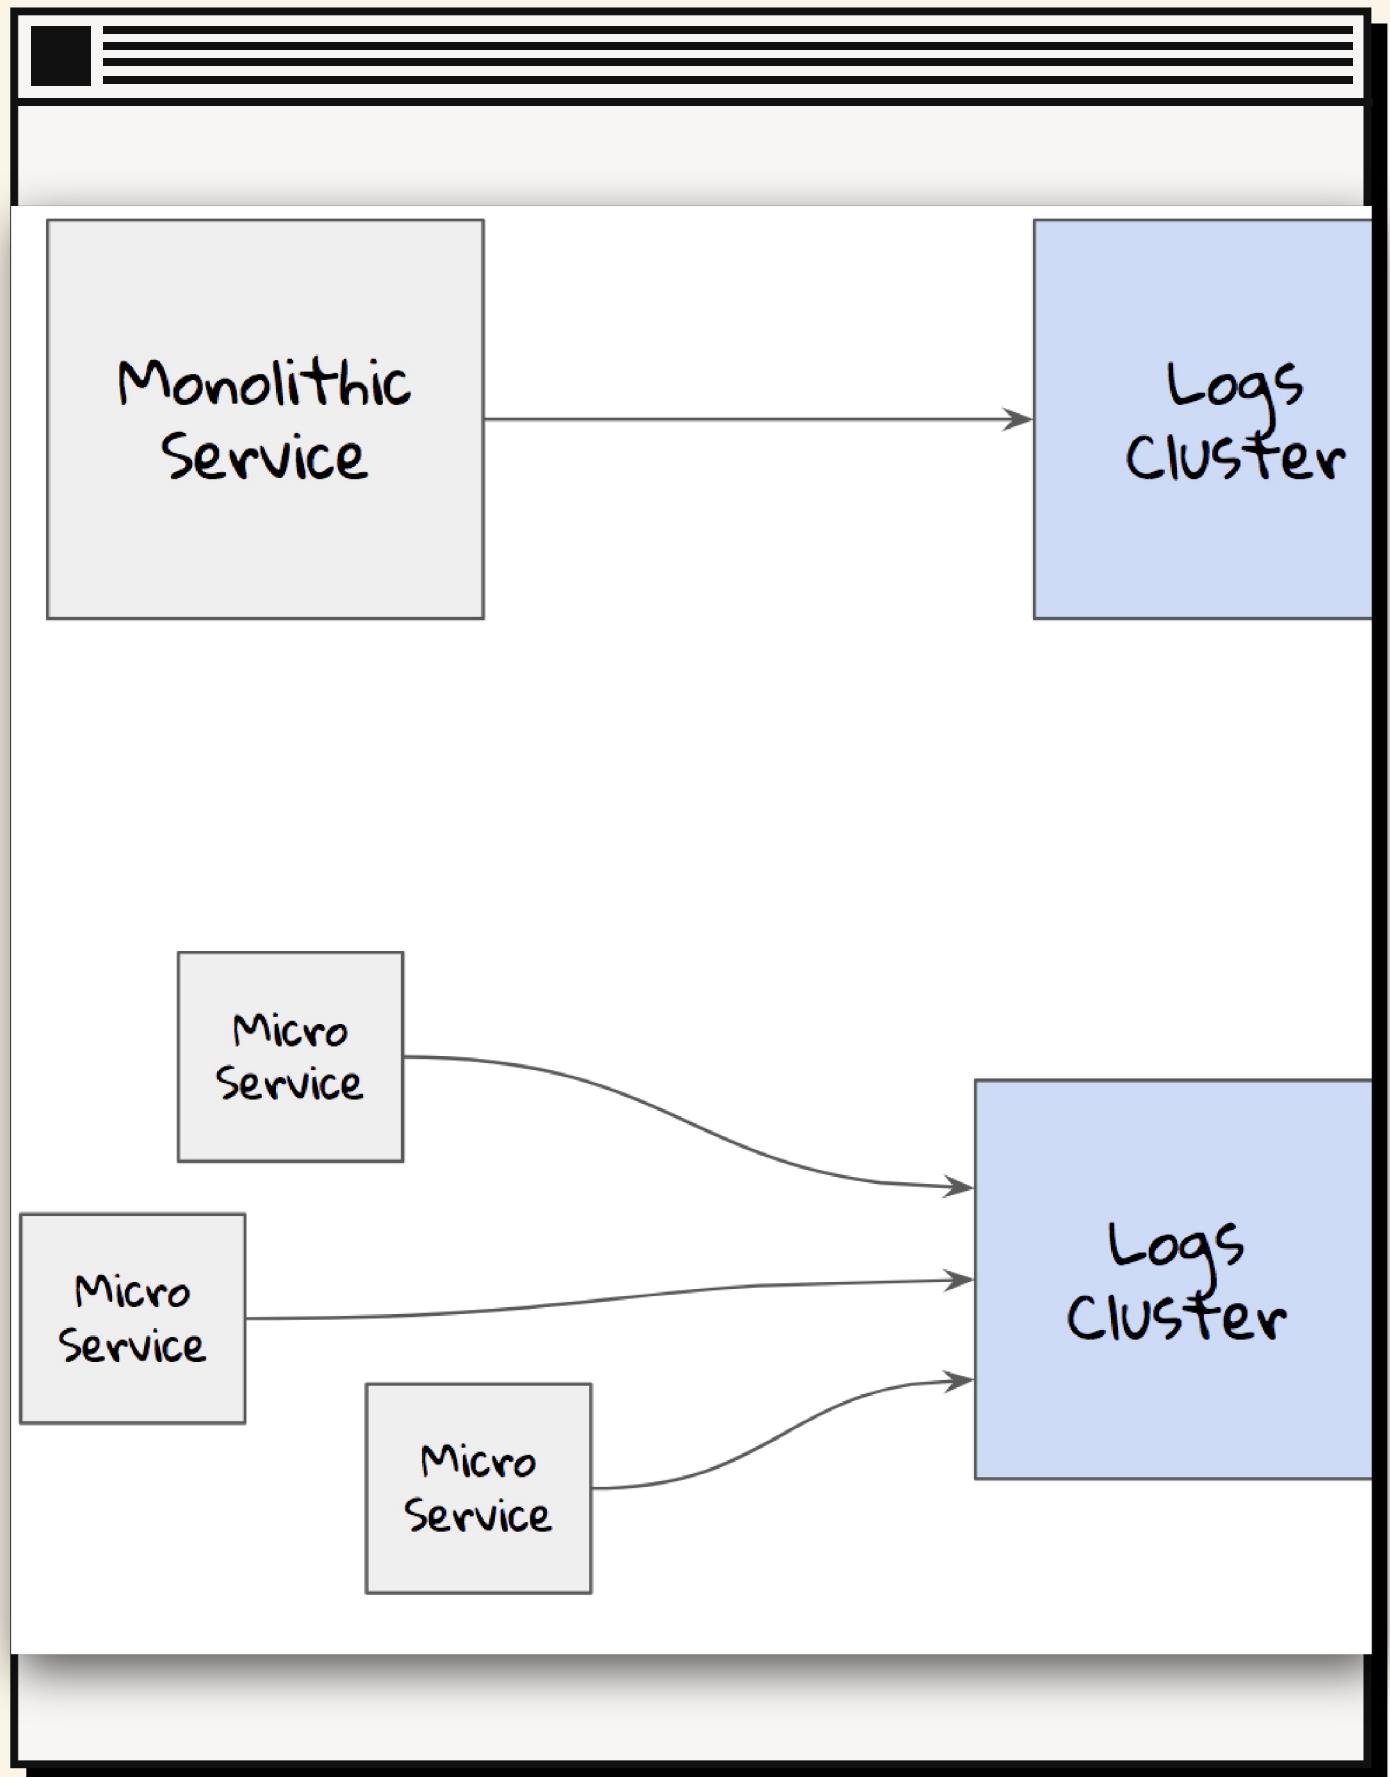
# Once upon a time

*There is a company who wants to adopt SRE practices.*

→ **Logging** - Tracing - Metrics **came along** to the picture

Good news: their  
architecture is quite  
simple

Bad news: the simplicity  
is not always



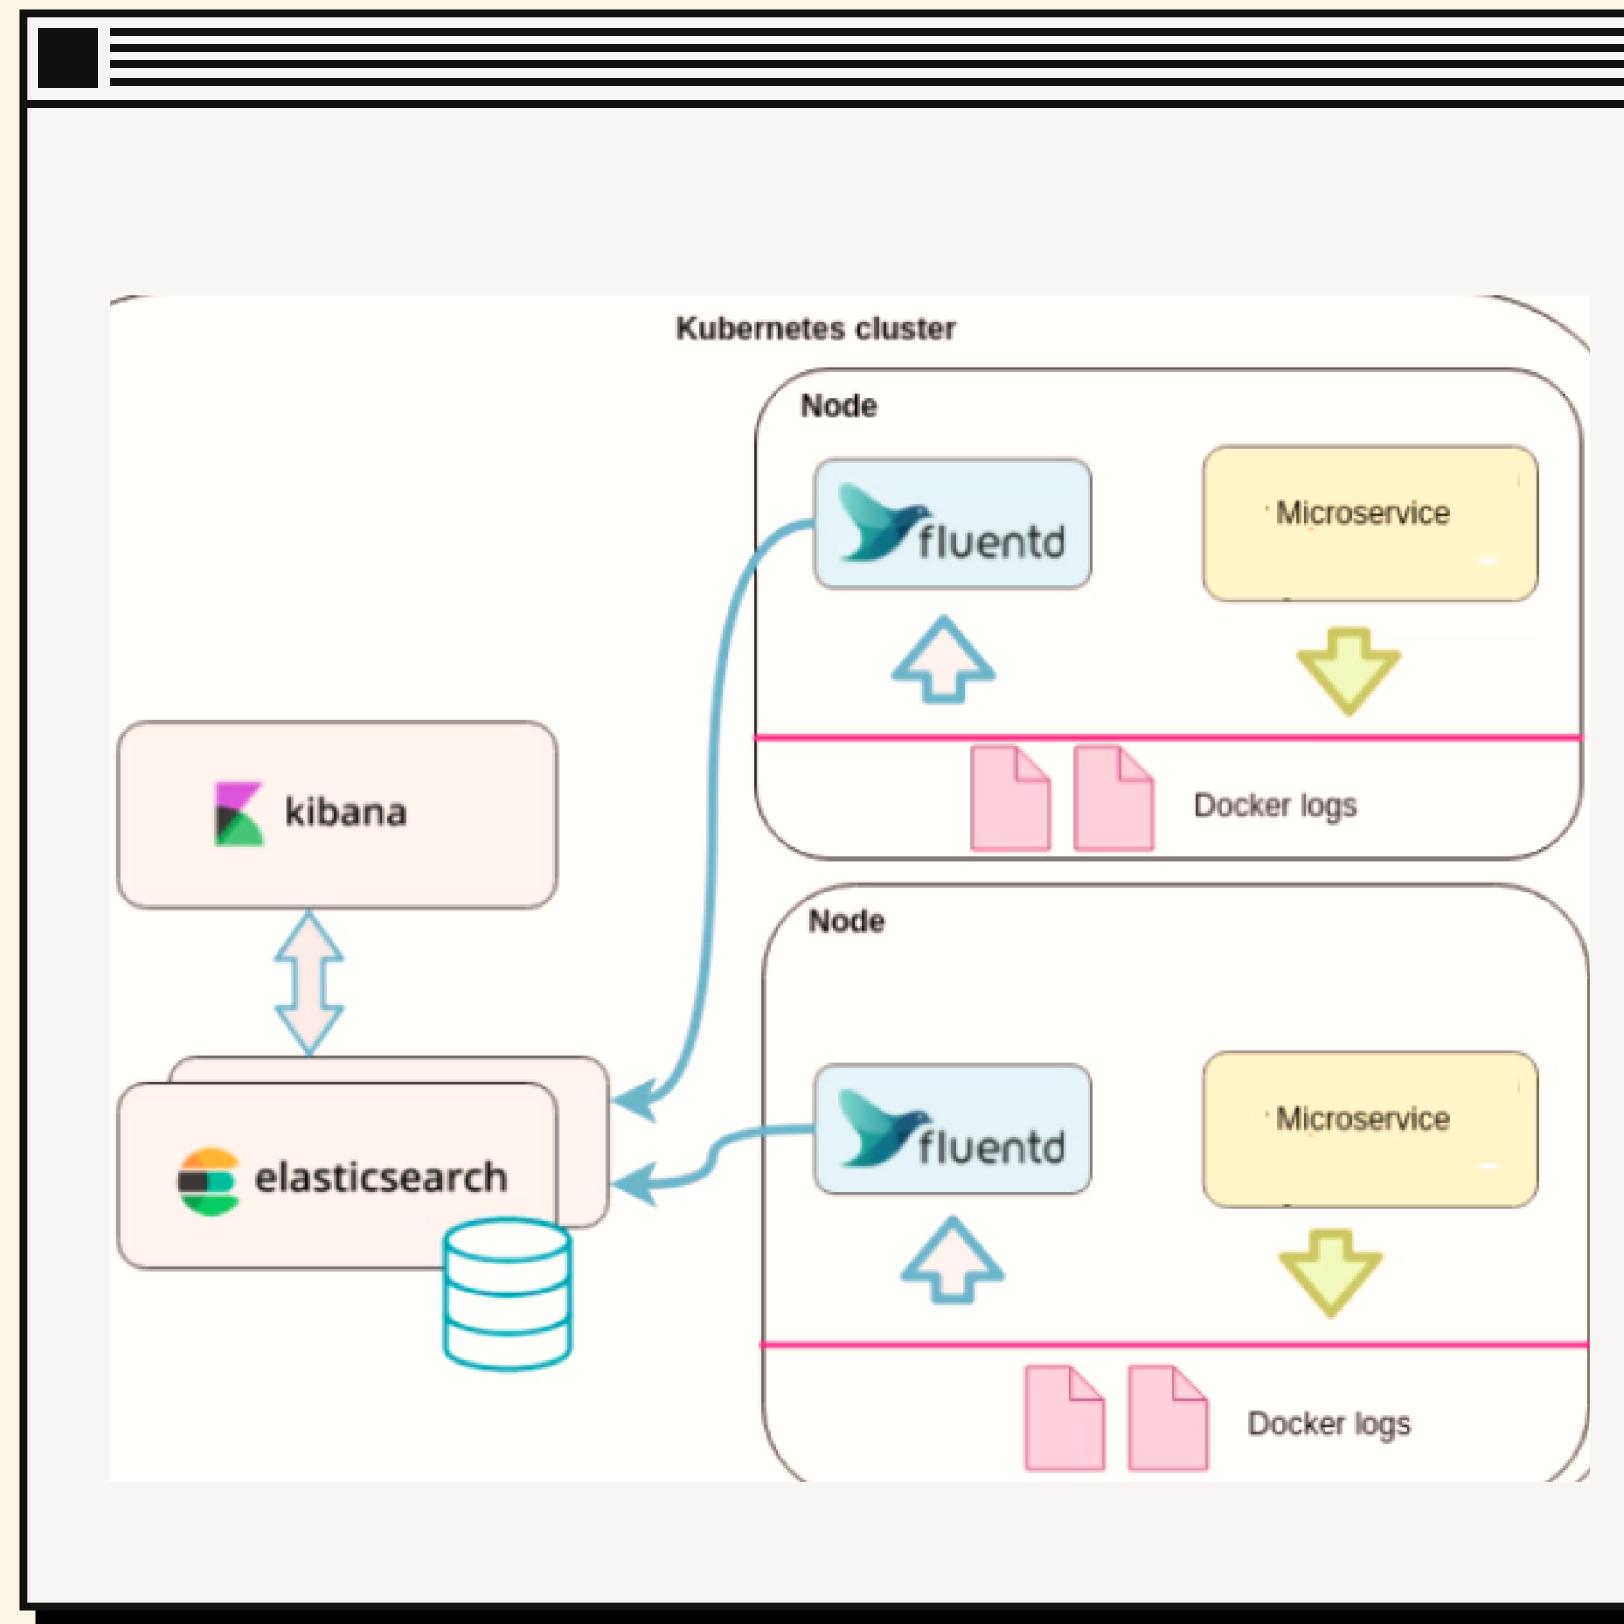


# And it's broken

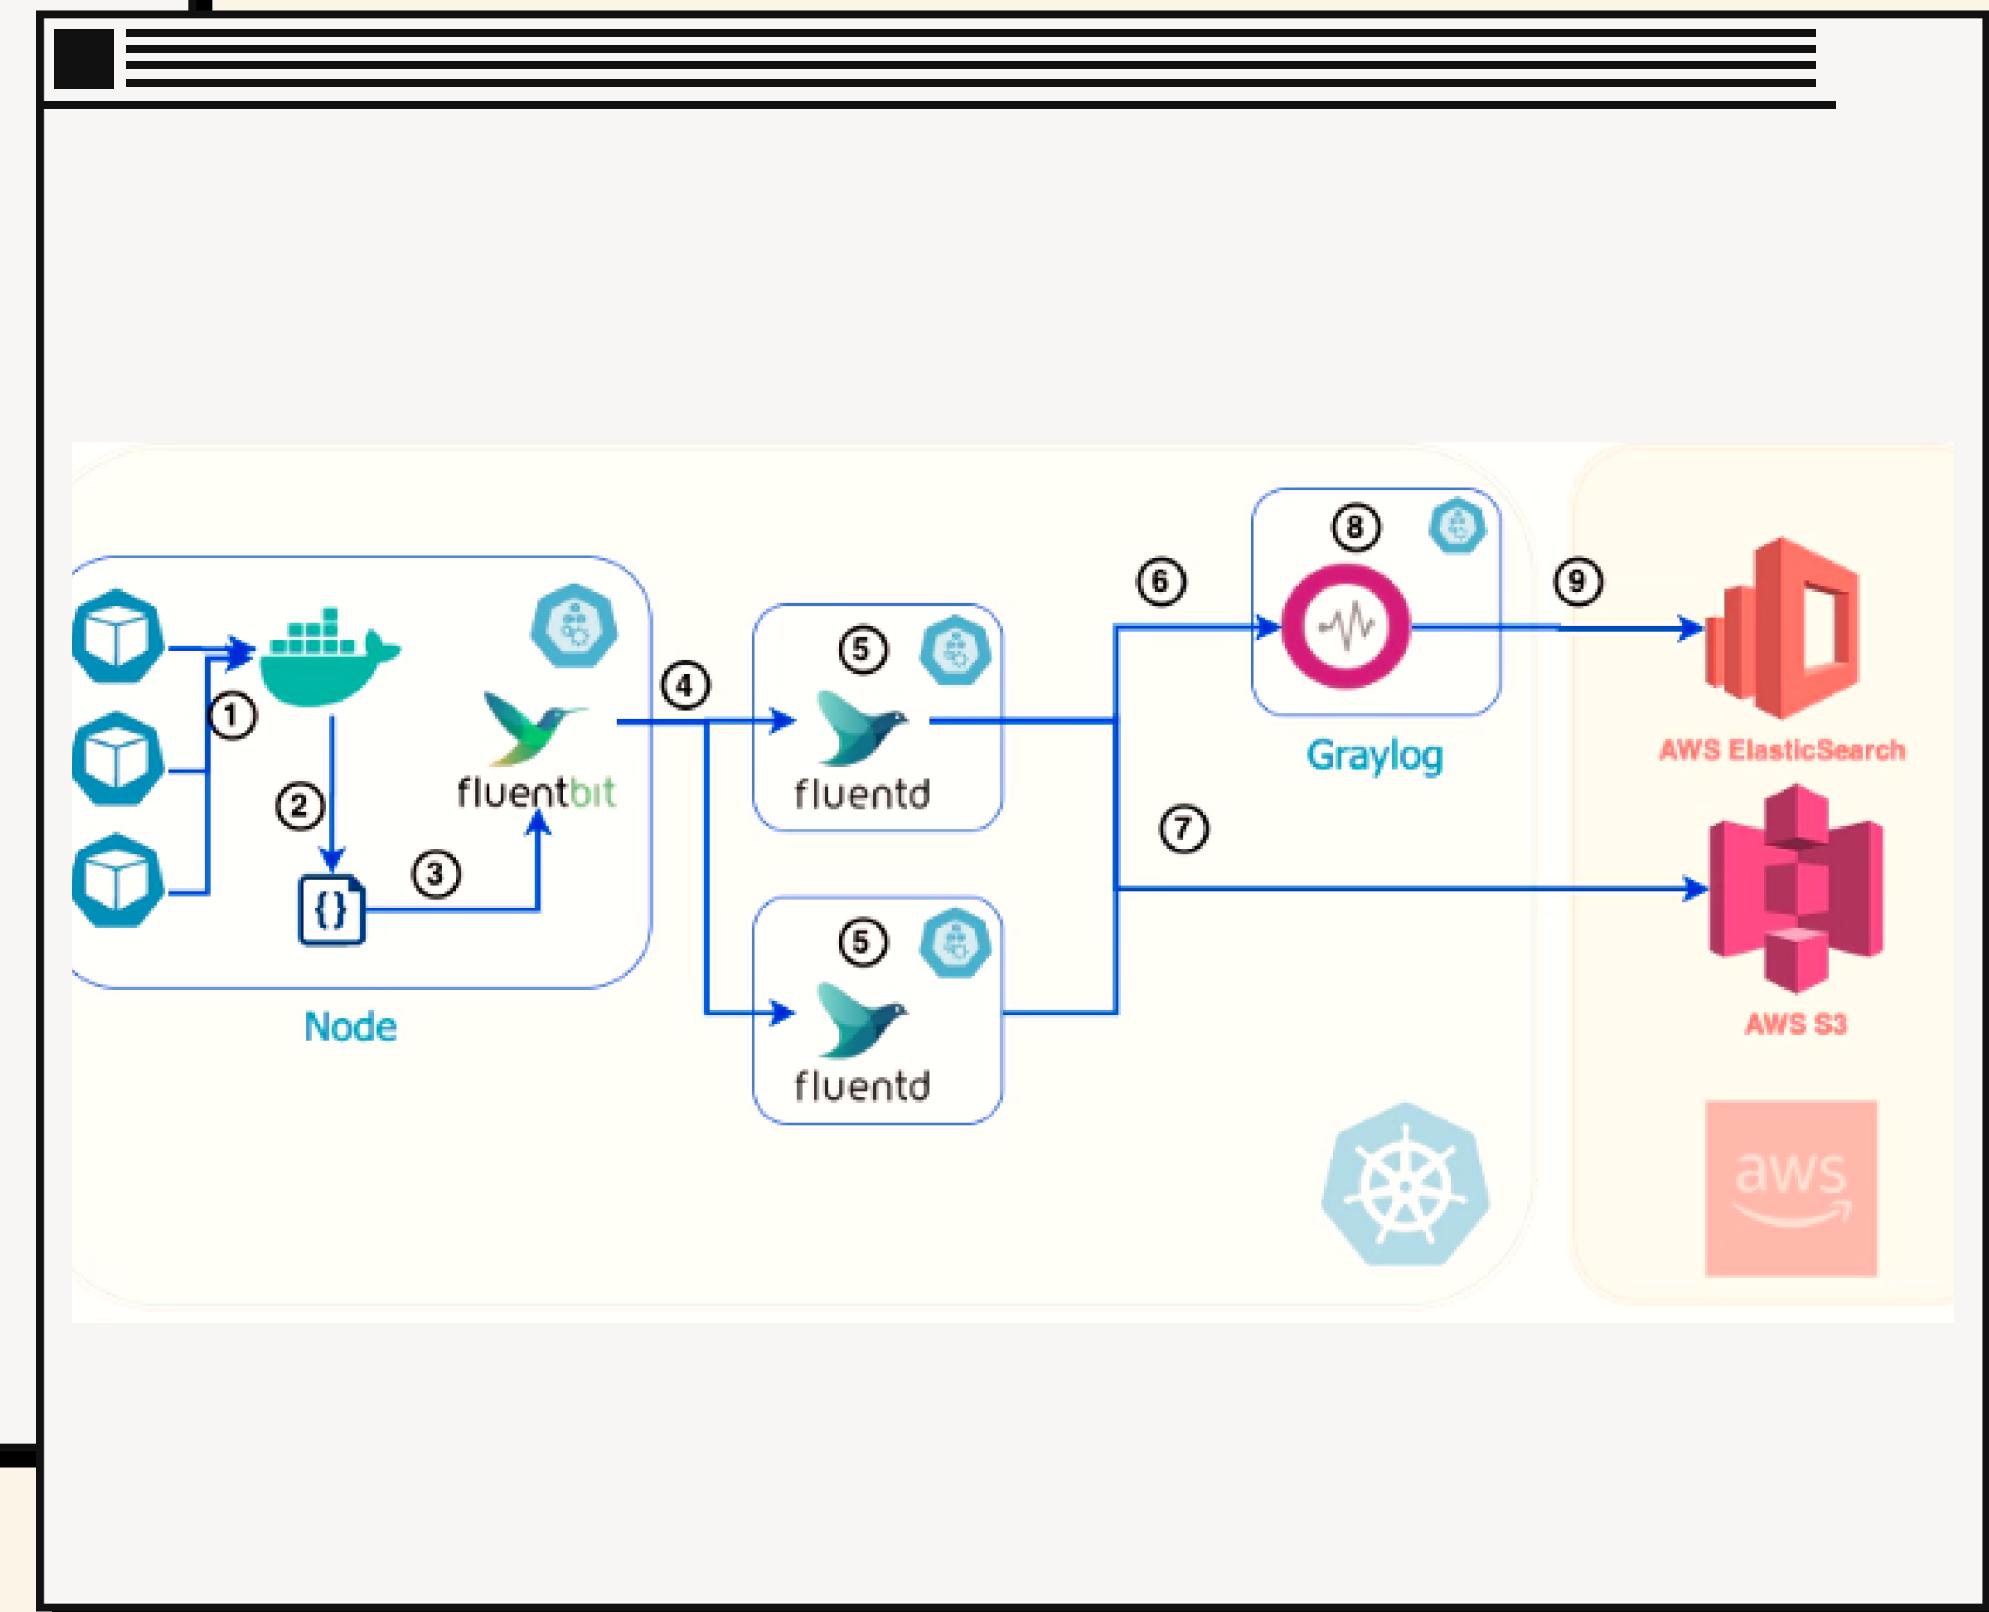
- It doesn't scale well
- So many unmeaningful events collected
- No one understands what is going on Logging Pipeline
- The rate of indexing failures are keep growing

- How to **normalise / standardise** all the log events ?
- How to **track all operations** which are happening in side your logging pipeline ?
- How to **distribute** the core configuration to every engineers ?
- How to **prevent** the applications who try to abuse your logging pipeline ?

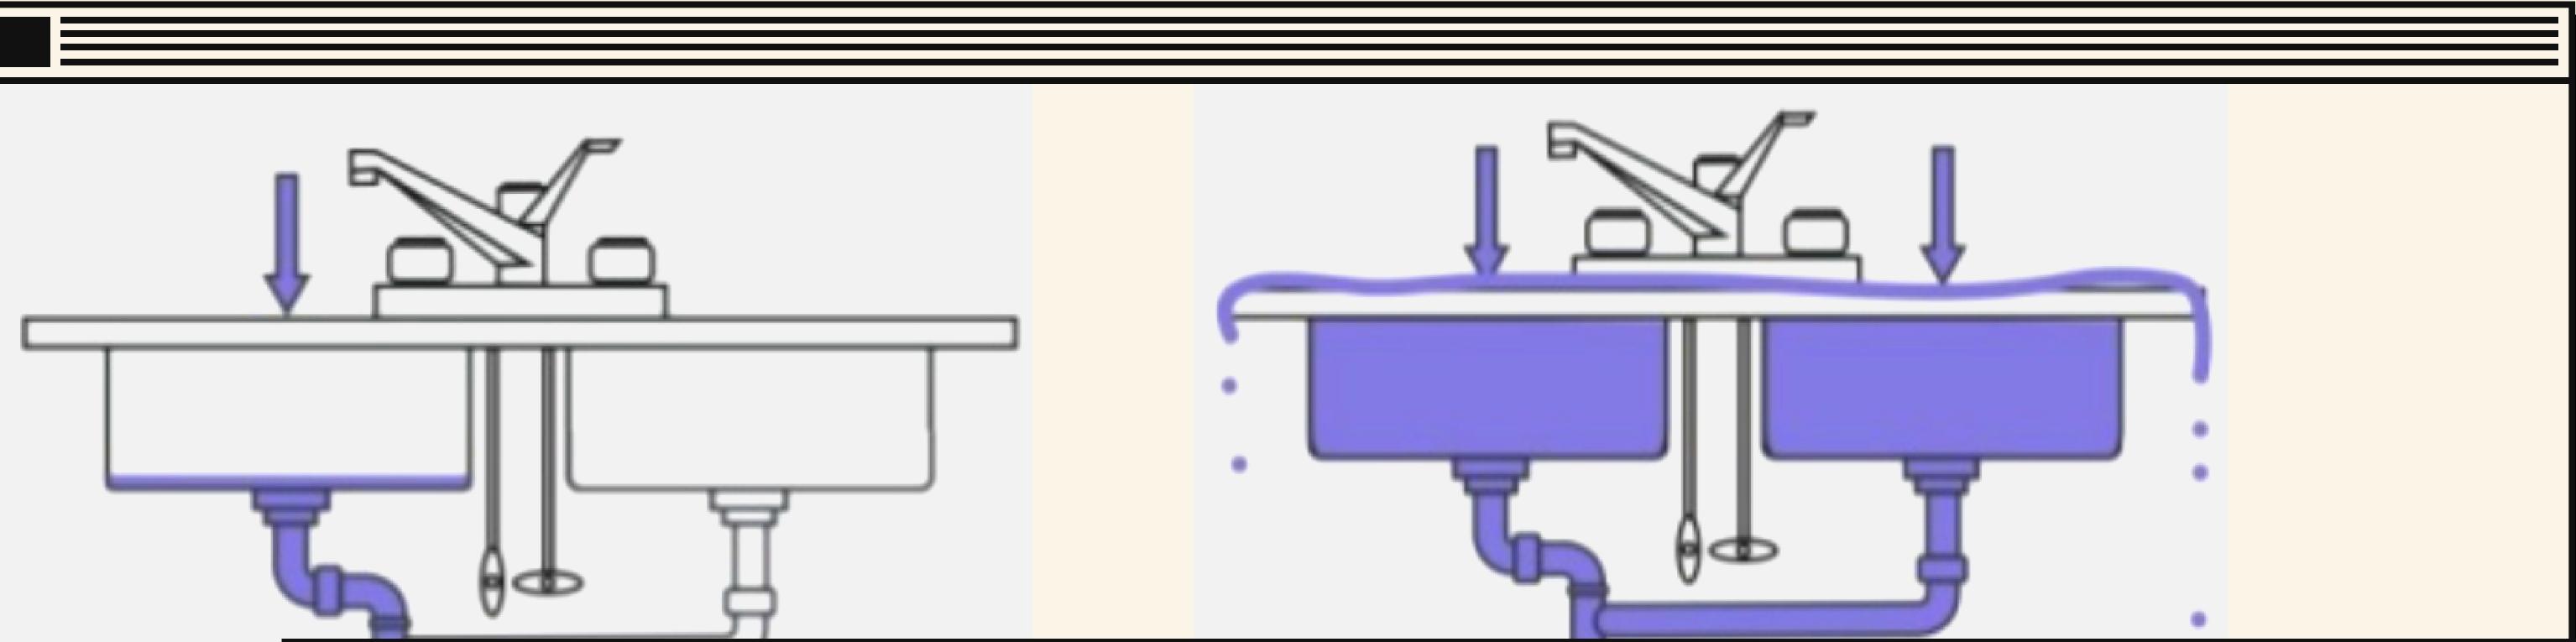
*Leverage from logging application to the platform is a challenge*



This is the old

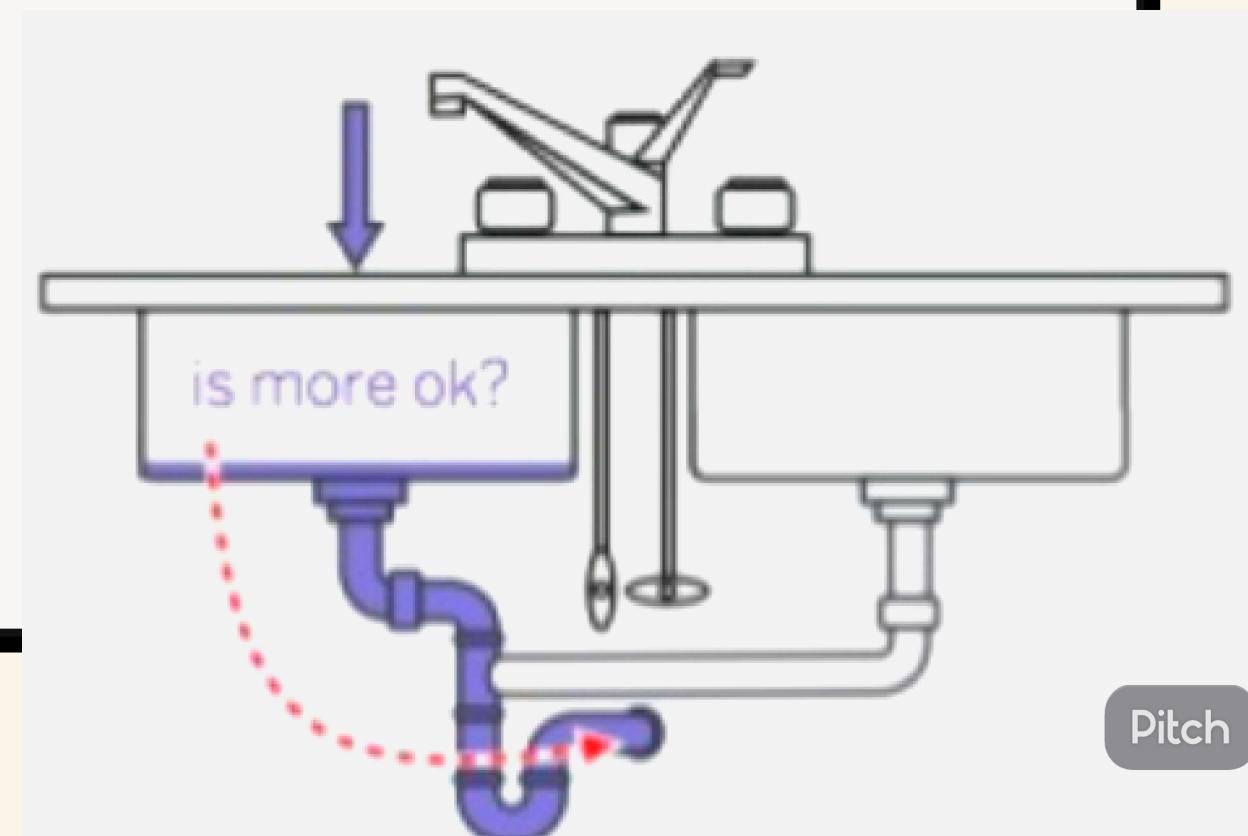


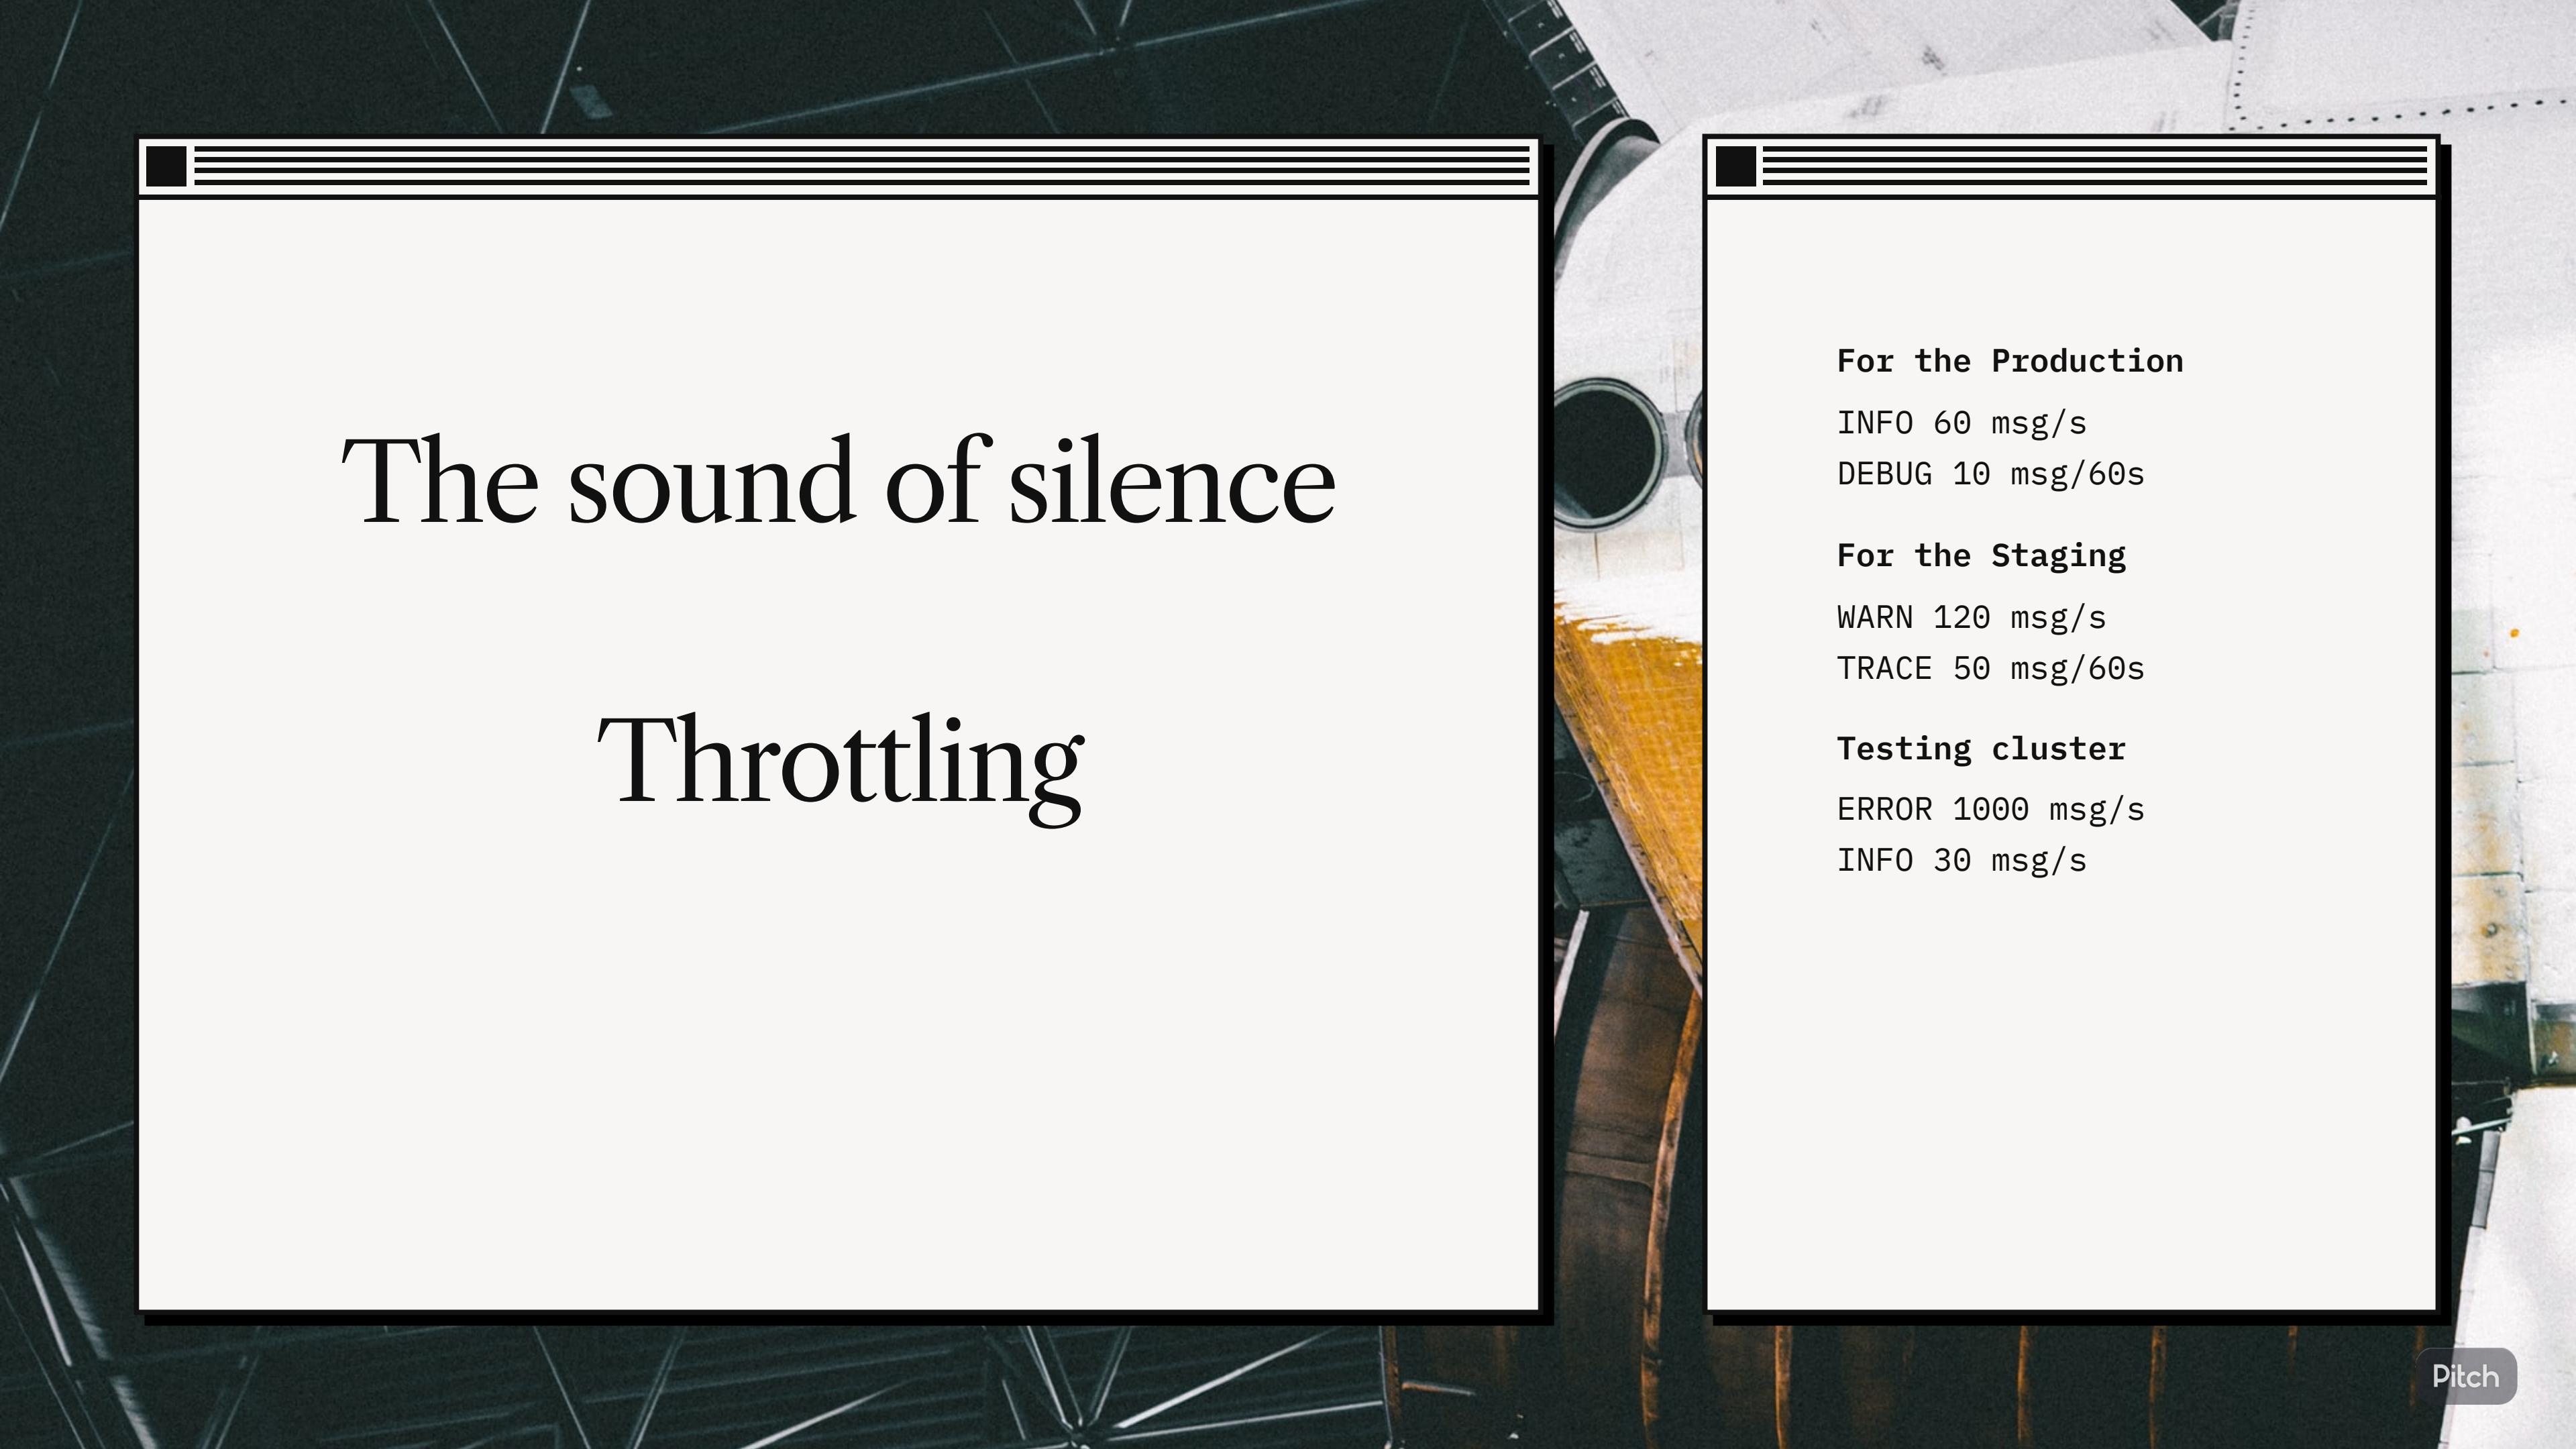
This is the new



### BackPressure mechanism

The main reason to use FluentBit is BackPressure mechanism which FluentD itself didn't have.





# The sound of silence

## Throttling

**For the Production**

INFO 60 msg/s  
DEBUG 10 msg/60s

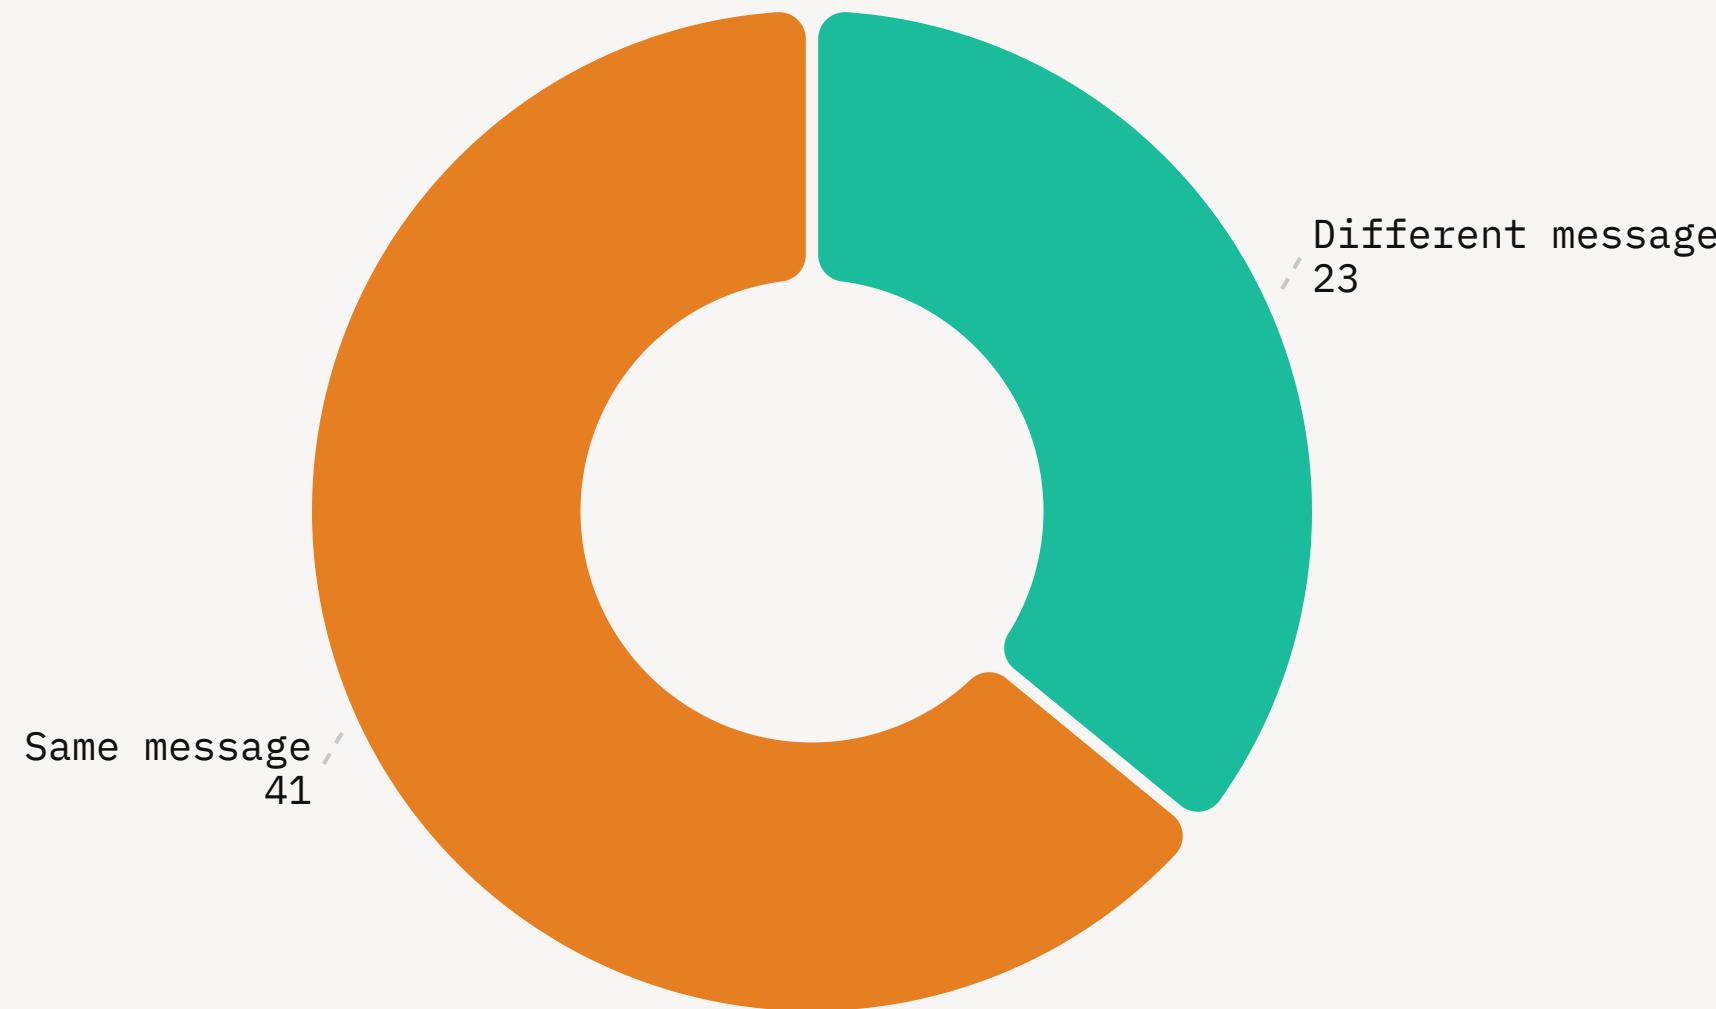
**For the Staging**

WARN 120 msg/s  
TRACE 50 msg/60s

**Testing cluster**

ERROR 1000 msg/s  
INFO 30 msg/s

# Sampling



## Mechanism

Suppress same messages

```
<match foo.**>
  @type      suppress
  interval   10
  num        2
  attr_keys  host,message
</match>
```

```
2020-11-22T11:22:33 foo.info
{"id":1,"host":"web01","message":"error!!"}
2020-11-22T11:22:34 foo.info
{"id":2,"host":"web01","message":"error!!"}
* 2020-11-22T11:22:35 foo.info
{"id":3,"host":"web01","message":"error!!"}
* 2020-11-22T11:22:36 foo.info
 {"id":4,"host":"web01","message":"error!!"}
2020-11-22T11:22:45 foo.info
 {"id":9,"host":"web01","message":"error!!"}
```

## TRACKING THE MESSAGE SIZE TO GET MORE VISIBILITY on STORAGE

*Integrate Flow Counter & Prometheus plugin to achieve this*



- *The weakness thing is the strongest point*

**FluentD Multiple Workers is ready to use**

```
<system>
  workers 10
</system>
```

## Distributed configuration

- FluentD introduces **include syntax**

```
<system>
@include teamA.conf
@include teamB.conf
@include teamC.conf
</system>
```

- Enable FluentD RPC

```
<system>
  rpc_endpoint
    127.0.0.1:24444
</system>
$ curl
  http://127.0.0.1:24444/
  api/plugins.flushBuffer
    s
    {"ok":true}
```

- Utilise ConfigMap & MountPoint

```
kubectl get cm
NAME          DATA   AGE
fluentd-teamA.conf  1     8d
fluentd-teamB.conf  1     8d
fluentd-teamC.conf  1     8d
```

- Engineers now have full visibility on the logging pipeline

## Log fields standardisation

Name	Format	Example	Description
Customer_Id	String	Alice-120	🐒 Monkey
Order_Id	Integer	12385910	🦏 Rhino
Trace_Id	Numeric	9292.14	🦩 Flamingo
Upstream_Host	String	<a href="http://checkout.teama.svc.cluster.local">http://checkout.teama.svc.cluster.local</a>	🦄 Unicorn

Document all fields info to allow developers get references easier.

# What if they don't follow ?

```
{"type":"mapper_parsing_exception","reason":"failed to parse field [duration] of type [float] in document with id 'da519231-1aaf-11eb-bb08-b22e566b37e4','caused_by":{"type":"number_format_exception","reason":"For input string: \"3.196501964s\"\"}}
```



ElasticSearch well known issue

# FluentBit filtering by



The filter will check each log field to see if it matches the data type in our list. If it matches, it will leave it as is.  
If it does not match, it will either just force/cast the value to a string value, or in case of a required numeric field it will try to parse it to a number

## But why GrayLog ?



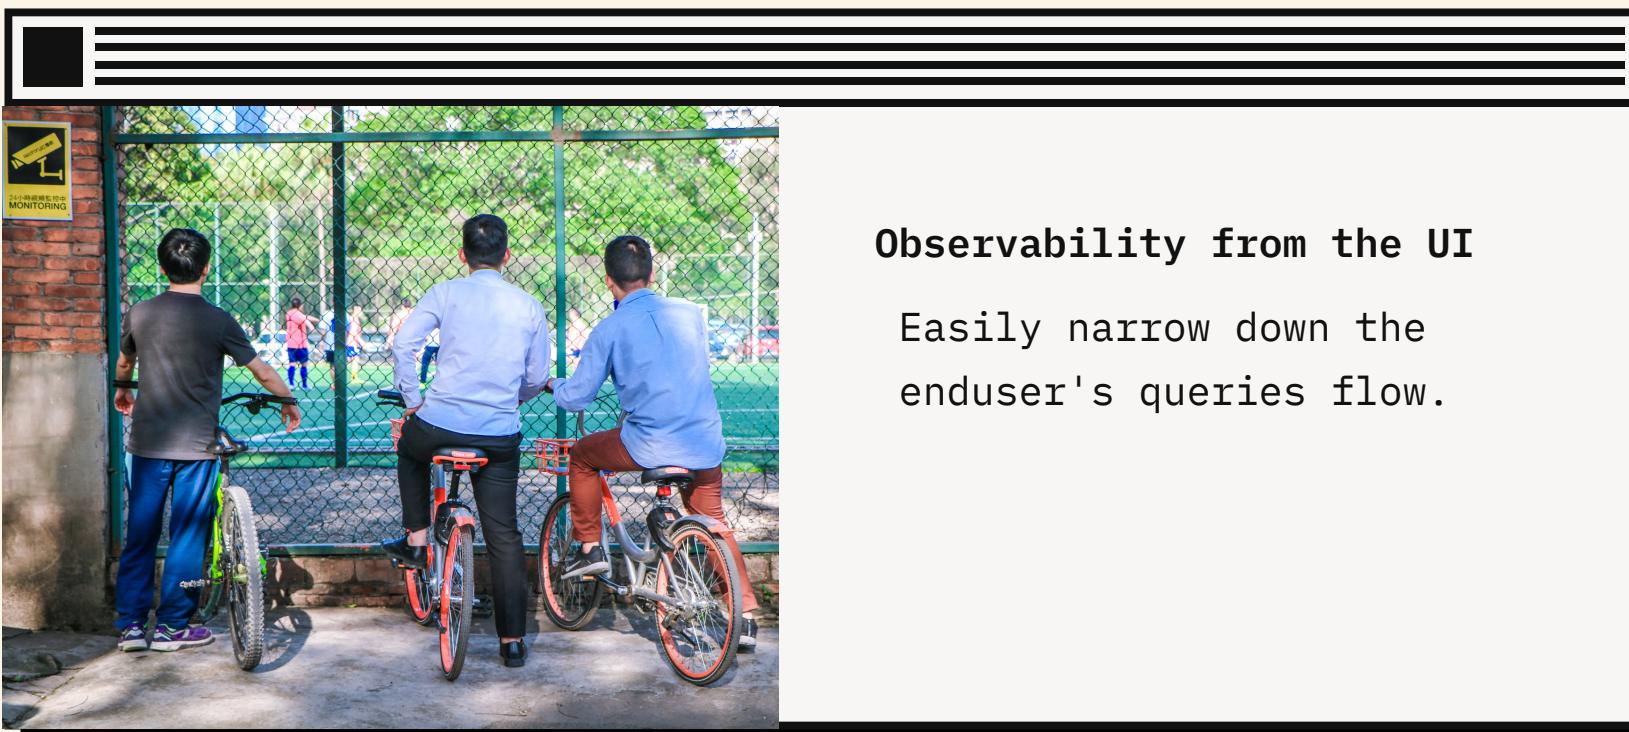
### Authorisation

To grant explicit permission  
for who wants to do what



### The abstraction layer for ES

To help us manage indices -  
index & other noisy things  
from ElasticSearch



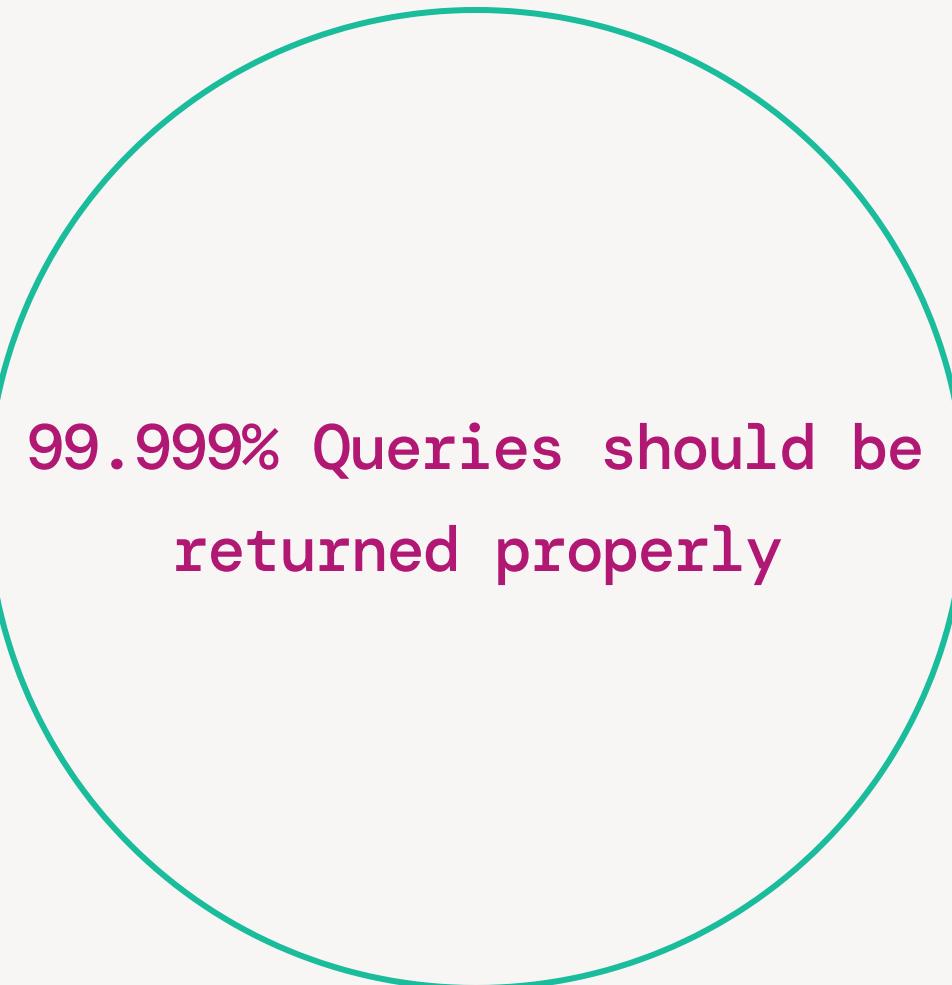
### Observability from the UI

Easily narrow down the  
enduser's queries flow.



### Manipulate data

To create the different  
pattern for what's exactly you  
want to show from the data.



From End User's perspective



From Developer's perspective

# LOG WELL

# LIVE WELL

